

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, r2_score

# Set plotting style
sns.set(style="whitegrid")

# 3. DATA LOADING & INITIAL EXPLORATION

# Read the CSV file (update the path as needed)
df = pd.read_csv("/content/drive/MyDrive/Machine Learning/COMP1816_Housing_Dataset_Regression.csv")

# Check for null values in each column and overall
null_counts = df.isnull().sum()
print("Null values per column:\n", null_counts)

total_nulls = df.isnull().sum().sum()
print(f"\nTotal number of null values: {total_nulls}")

# Display rows with any null values
print("\nRows with null values:")
print(df[df.isnull().any(axis=1)])

# Drop rows with null values and display dataframe info
df.dropna(inplace=True)
df.info()

# 4. DATA SPLITTING

# Separate features and target variable
X = df.drop(['median_house_value'], axis=1)
y = df['median_house_value']

# Split data into train and test sets (using last 190 samples for testing)
X_train, X_test = X.iloc[:-190], X.iloc[-190:]
y_train, y_test = y.iloc[:-190], y.iloc[-190:]

# Combine X and y for training set for easier preprocessing
train_data = pd.concat([X_train, y_train], axis=1)

# 5. EXPLORATORY DATA ANALYSIS (EDA)

# Plot histogram for training data
plt.figure(figsize=(15, 9))
train_data.hist(bins=50, figsize=(15, 9))
plt.tight_layout()
plt.show()

# Calculate and plot correlation heatmap for numerical features
numerical_features = train_data.select_dtypes(include=['number'])
correlation_matrix = numerical_features.corr()

plt.figure(figsize=(15, 9))
sns.heatmap(correlation_matrix, annot=True, cmap='YlGnBu')
plt.title("Correlation Heatmap (Training Data)")
plt.show()

# Scatter plot for geographic distribution colored by median house value
plt.figure(figsize=(15, 9))
sns.scatterplot(x='longitude', y='latitude', hue='median_house_value', data=train_data, palette='coolwarm')
plt.title("Geographic Scatter Plot")
plt.show()

# 6. DATA PREPROCESSING & FEATURE ENGINEERING

# --- Preprocessing for Training Data ---
# Log-transform selected numerical features (add 1 to avoid log(0))
for col in ['total_rooms', 'total_bedrooms', 'population', 'households']:
    train_data[col] = np.log(train_data[col] + 1)

```

```
# Check the distribution of 'ocean_proximity'
print("\nOcean Proximity Value Counts (Training):")
print(train_data['ocean_proximity'].value_counts())

# Convert categorical feature 'ocean_proximity' to dummy variables
train_data = train_data.join(pd.get_dummies(train_data.ocean_proximity)).drop(['ocean_proximity'], axis=1)

# Create new features
train_data['bedroom_ratio'] = train_data['total_bedrooms'] / train_data['total_rooms']
train_data['household_rooms'] = train_data['total_rooms'] / train_data['households']

# Plot updated correlation heatmap
plt.figure(figsize=(15, 9))
sns.heatmap(train_data.corr(), annot=True, cmap='YlGnBu')
plt.title("Correlation Heatmap with Engineered Features (Training Data)")
plt.show()
```

7. PREPARE DATA FOR MODELING

```
# Separate features and target from training data
X_train_proc = train_data.drop(['median_house_value'], axis=1)
y_train_proc = train_data['median_house_value']
```

```
# Standardize the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_proc)
```

8. TRAIN A LINEAR REGRESSION MODEL

```
reg = LinearRegression()
reg.fit(X_train_scaled, y_train_proc)
```

9. PREPROCESS THE TEST DATA

```
# Combine test features and target for preprocessing
test_data = X_test.join(y_test)
```

```
# Apply log transformation to numerical features
for col in ['total_rooms', 'total_bedrooms', 'population', 'households']:
    test_data[col] = np.log(test_data[col] + 1)
```

```
# Convert categorical 'ocean_proximity' to dummy variables
test_data = test_data.join(pd.get_dummies(test_data.ocean_proximity)).drop(['ocean_proximity'], axis=1)
```

```
# Create new features for test data
test_data['bedroom_ratio'] = test_data['total_bedrooms'] / test_data['total_rooms']
test_data['household_rooms'] = test_data['total_rooms'] / test_data['households']
```

```
# Ensure the test data has the same feature columns as training data
X_test_proc = test_data.reindex(columns=X_train_proc.columns, fill_value=0)
y_test_proc = test_data['median_house_value']
```

```
# Scale test features using the same scaler as training data
X_test_scaled = scaler.transform(X_test_proc)
```

10. EVALUATE THE LINEAR REGRESSION MODEL

```
y_pred_lr = reg.predict(X_test_scaled)
mse_lr = mean_squared_error(y_test_proc, y_pred_lr)
rmse_lr = np.sqrt(mse_lr)
r2_lr = r2_score(y_test_proc, y_pred_lr)

print("\nLinear Regression Evaluation:")
print(f"Mean Squared Error (MSE): {mse_lr:.2f}")
print(f"Root Mean Squared Error (RMSE): {rmse_lr:.2f}")
print(f"R-squared (R²): {r2_lr:.2f}")
```

11. TRAIN A RANDOM FOREST MODEL (Directly on Test Data for Demo)

```
forest = RandomForestRegressor()
forest.fit(X_test_scaled, y_test_proc)
```

12. GRID SEARCH FOR RANDOM FOREST HYPERPARAMETERS

```
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_features': [2, 4, 6, 8],
    'min_samples_split': [2, 4],
```

```

    'max_depth': [None, 4, 8]
}

grid_search = GridSearchCV(RandomForestRegressor(), param_grid, cv=5,
                           scoring="neg_mean_squared_error", return_train_score=True)
grid_search.fit(X_train_scaled, y_train_proc)

best_rf = grid_search.best_estimator_
print("\nGrid Search Best Estimator (Random Forest):")
print(best_rf)
print(f"Test Score (R²): {best_rf.score(X_test_scaled, y_test_proc):.2f}")

```

13. MODEL COMPARISON: DECISION TREE, RANDOM FOREST, GRADIENT BOOSTING

```

# Define models
models = {
    "Decision Tree": DecisionTreeRegressor(random_state=42),
    "Random Forest": RandomForestRegressor(n_estimators=100, random_state=42),
    "Gradient Boosting": GradientBoostingRegressor(n_estimators=100, random_state=42)
}

results = {}

# Evaluate each model
for name, model in models.items():
    model.fit(X_train_scaled, y_train_proc)
    y_pred = model.predict(X_test_scaled)
    results[name] = {
        "MSE": mean_squared_error(y_test_proc, y_pred),
        "RMSE": np.sqrt(mean_squared_error(y_test_proc, y_pred)),
        "R²": r2_score(y_test_proc, y_pred)
    }

print("\nModel Comparison:")
print("{:<20} {:<10} {:<10} {:<10}".format('Model', 'MSE', 'RMSE', 'R²'))
for model_name, metrics in results.items():
    print("{:<20} {:<10.2f} {:<10.2f} {:<10.2f}".format(
        model_name, metrics['MSE'], metrics['RMSE'], metrics['R²']
    ))

# Print Linear Regression results for reference
print("\nLinear Regression Reference:")
print("{:<20} {:<10.2f} {:<10.2f} {:<10.2f}".format(
    "Linear Regression", mse_lr, rmse_lr, r2_lr
))

```

14. FINAL HYPERPARAMETER TUNING & EVALUATION WITH RANDOM FOREST

```

# Defining a simpler parameter grid for final tuning
param_grid_final = {
    'n_estimators': [100, 200],
    'max_depth': [None, 10],
    'min_samples_split': [2, 5]
}

grid_search_final = GridSearchCV(
    estimator=RandomForestRegressor(random_state=42),
    param_grid=param_grid_final,
    cv=5,
    scoring='neg_mean_squared_error',
    n_jobs=-1
)

grid_search_final.fit(X_train_scaled, y_train_proc)
best_model = grid_search_final.best_estimator_

print("\nFinal Grid Search Best Parameters:")
print(grid_search_final.best_params_)
print(f"Best RMSE (Training CV): {np.sqrt(-grid_search_final.best_score_):.2f}")

# Feature Importance
importances = best_model.feature_importances_
feature_names = X_train_proc.columns
sorted_idx = np.argsort(importances)[::-1]

plt.figure(figsize=(12, 6))
plt.bar(range(len(feature_names)), importances[sorted_idx], align='center')
plt.xticks(range(len(feature_names)), feature_names[sorted_idx], rotation=90)
plt.title("Random Forest Feature Importances")
plt.xlabel("Features")
plt.ylabel("Importance Score")

```

```
plt.tight_layout()
plt.show()

#Residual Analysis
y_pred_best = best_model.predict(X_test_scaled)
residuals = y_test_proc - y_pred_best

plt.figure(figsize=(10, 6))
plt.scatter(y_pred_best, residuals, alpha=0.5)
plt.hlines(y=0, xmin=y_pred_best.min(), xmax=y_pred_best.max(), colors='red', linestyle='--')
plt.title("Residual Plot for Best Random Forest Model")
plt.xlabel("Predicted Values")
plt.ylabel("Residuals (Actual - Predicted)")
plt.grid(True)
plt.show()

#Final Evaluation
mse_final = mean_squared_error(y_test_proc, y_pred_best)
rmse_final = np.sqrt(mse_final)
r2_final = best_model.score(X_test_scaled, y_test_proc)

print("\nFinal Evaluation Metrics (Best Random Forest):")
print(f"MSE: {mse_final:.2f}")
print(f"RMSE: {rmse_final:.2f}")
print(f"R2 Score: {r2_final:.2f}")
```

Null values per column:

```
No.      0
longitude 0
latitude  0
housing_median_age 0
total_rooms 0
total_bedrooms 12
population 0
households 0
median_income 0
median_house_value 0
ocean_proximity 7
dtype: int64
```

Total number of null values: 19

Rows with null values:

No.	longitude	latitude	housing_median_age	total_rooms	\
72	73	-122.08	37.88	26	2947
93	94	-119.80	36.75	46	2625
98	99	-119.82	36.81	25	3305
168	169	-118.28	34.25	29	2559
236	237	-118.38	34.05	49	702
548	549	-117.87	33.83	27	2287
585	586	-121.26	38.74	22	7173
595	596	-121.04	39.00	21	4059
603	94	-119.80	36.75	46	2625
608	99	-119.82	36.81	25	3305
621	622	-116.21	33.75	22	894
740	741	-117.02	32.66	19	771
786	787	-122.45	37.77	52	2602
792	793	-122.50	37.75	45	1620
821	822	-122.39	37.59	32	4497
893	894	-121.94	36.97	31	1738
981	982	-119.06	34.24	21	7436
988	989	-118.69	34.18	11	1177
992	993	-121.52	38.57	43	2360

	total_bedrooms	population	households	median_income	\
72	NaN	825	626	2.9330	
93	593.0	1368	551	1.5273	
98	NaN	1149	500	5.0698	
168	NaN	1886	769	2.6036	
236	NaN	458	187	4.8958	
548	NaN	1140	351	5.6163	
585	1314.0	3526	1316	3.3941	
595	730.0	1874	693	4.8051	
603	593.0	1368	551	1.5273	
608	NaN	1149	500	5.0698	
621	NaN	830	202	3.0673	
740	NaN	376	108	6.6272	
786	NaN	1330	647	3.5435	
792	NaN	941	328	4.3859	
821	NaN	1846	715	6.1323	
893	422.0	746	355	2.5172	
981	984.0	2982	988	7.6775	
988	NaN	415	119	10.0472	
992	471.0	1041	452	2.8900	

	median_house_value	ocean_proximity
72	85000	NEAR BAY
93	59000	NaN
98	150900	INLAND
168	162100	<1H OCEAN
236	333600	<1H OCEAN
548	231000	<1H OCEAN
585	135900	NaN
595	174300	NaN
603	59000	NaN
608	150900	INLAND
621	68200	INLAND
740	273600	NEAR OCEAN
786	278600	NEAR BAY
792	270200	NEAR OCEAN
821	500001	NEAR OCEAN
893	330800	NaN
981	391200	NaN
988	500001	<1H OCEAN
992	86200	NaN

<class 'pandas.core.frame.DataFrame'>

Index: 981 entries, 0 to 999

Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
0	No.	981 non-null	int64