

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# prompt: /content/drive/MyDrive/COMP1816_Titanic_Dataset_Classification.csv this is my file read the csv file with python
```

```
import pandas as pd
# Load the dataset
df = pd.read_csv('/content/drive/MyDrive/COMP1816_Titanic_Dataset_Classification.csv')
```

```
# Display the first few rows of the dataframe
print(df.head())
```


```

PassengerId  Pclass  Name  Sex
0            1     3.0  Braund, Mr. Owen Harris  male
1            2     1.0           NaN  female
2            3     3.0  Heikkinen, Miss. Laina  female
3            4     1.0  Futrelle, Mrs. Jacques Heath (Lily May Peel)  female
4            5     3.0  Allen, Mr. William Henry  male

Age  SibSp  Parch  Ticket No.  Fare  Embarked  Survival
0  22.0    1.0    0.0    A/5 21171   7.2500      S      0.0
1  38.0    1.0    0.0    PC 17599  71.2833      C      1.0
2  26.0    0.0    0.0  STON/O2. 3101282   7.9250      S      1.0
3  35.0    1.0    0.0    113803  53.1000      S      1.0
4  35.0    0.0    0.0    373450   8.0500      S      0.0
```


```
sns.countplot(x='Survival', data= df, hue ='Pclass')
```

```
<Axes: xlabel='Survival', ylabel='count'>
```



so here it is clear that majority of the people did not survived from the ticketclass 3

```
sns.displot(df['Age'], kde=False)
plt.xlim(0, 100) # Set the x-axis limits
plt.show()
```




```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 890 entries, 0 to 889
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   PassengerId  890 non-null    int64
1   Pclass      880 non-null    float64
2   Name        885 non-null    object
3   Sex         890 non-null    object
4   Age         715 non-null    float64
5   SibSp       888 non-null    float64
6   Parch       888 non-null    float64
7   Ticket No.   888 non-null    object
8   Fare        888 non-null    float64
9   Embarked    884 non-null    object
10  Survival     886 non-null    float64
dtypes: float64(6), int64(1), object(4)
memory usage: 76.6+ KB
```

```
df.isnull().sum()
```

```
0
PassengerId  0
Pclass      10
Name         5
Sex          0
Age        175
SibSp        2
Parch        2
Ticket No.   2
Fare         2
Embarked     6
Survival     4
```

```
sns.boxplot(x='Pclass', y='Age', data=df)
plt.ylim(0, 100) # Set the y-axis limits to 0-100
plt.show()
```



```
# Calculate the mean age for each passenger class
mean_age_by_pclass = df.groupby('Pclass')['Age'].mean()

mean_age_by_pclass
```

	Age
Pclass	
1.0	38.149022
2.0	29.793118
3.0	33.489801

```
print(df[df['Pclass']==1]['Age'].mean())
print(df[df['Pclass']==2]['Age'].mean())
print(df[df['Pclass']==3]['Age'].mean())
```

```
38.14902173913043
29.793117647058825
33.48980056980057
```

```
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
```

```
def fill_in_na_values(cols):
    age = cols[0]
    pclass = cols[1]

    if pd.isnull(age):
        if pclass == 1:
            return round(df[df['Pclass']==1]['Age'].mean())
        elif pclass == 2:
            return round(df[df['Pclass']==2]['Age'].mean())
        else:
            return round(df[df['Pclass']==3]['Age'].mean())
    else:
        return age

df['Age']=df[['Age', 'Pclass']].apply(fill_in_na_values, axis=1)
```

```
<ipython-input-24-fe67a078db9b>:2: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a futu
age = cols[0]
<ipython-input-24-fe67a078db9b>:3: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a futu
pclass = cols[1]
```

```
# Fill missing Pclass values (if any)
df['Pclass'].fillna(df['Pclass'].mode()[0], inplace=True)

# Function to fill missing Age values
def fill_in_na_values(cols, df):
    age, pclass = cols
    if pd.isnull(age):
        return round(df[df['Pclass'] == pclass]['Age'].mean())
    return age

# Apply function to fill missing Age values
df['Age'] = df[['Age', 'Pclass']].apply(lambda row: fill_in_na_values(row, df), axis=1)

print(df)
```

	PassengerId	Pclass	Name \
0	1	3.0	Braund, Mr. Owen Harris
1	2	1.0	NaN
2	3	3.0	Heikkinen, Miss. Laina
3	4	1.0	Futrelle, Mrs. Jacques Heath (Lily May Peel)
4	5	3.0	Allen, Mr. William Henry
..	...	...	...
885	886	2.0	Montvila, Rev. Juozas
886	887	1.0	Graham, Miss. Margaret Edith
887	888	3.0	Johnston, Miss. Catherine Helen "Carrie"
888	889	1.0	Behr, Mr. Karl Howell
889	890	3.0	Dooley, Mr. Patrick

	Sex	Age	SibSp	Parch	Ticket No.	Fare	Embarked	Survival
0	male	22.0	1.0	0.0	A/5 21171	7.2500	S	0.0
1	female	38.0	1.0	0.0	PC 17599	71.2833	C	1.0
2	female	26.0	0.0	0.0	STON/O2. 3101282	7.9250	S	1.0
3	female	35.0	1.0	0.0	113803	53.1000	S	1.0
4	male	35.0	0.0	0.0	373450	8.0500	S	0.0
..	...	...	...	...	...	...	...	...
885	male	27.0	0.0	0.0	211536	13.0000	S	0.0
886	female	19.0	0.0	0.0	112053	30.0000	S	1.0
887	female	33.0	1.0	2.0	W./C. 6607	23.4500	S	0.0
888	male	26.0	0.0	0.0	111369	30.0000	C	1.0
889	male	32.0	0.0	0.0	370376	7.7500	Q	0.0

[890 rows x 11 columns]

<ipython-input-25-8fcf2ce5818d>:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through  
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[c

```
df['Pclass'].fillna(df['Pclass'].mode()[0], inplace=True)
```

```
df.isnull().sum()
```

	0
PassengerId	0
Pclass	0
Name	5
Sex	0
Age	0
SibSp	2
Parch	2
Ticket No.	2
Fare	2
Embarked	6
Survival	4

```
def fill_in_na_values(cols):
    age = cols[0]
    pclass = cols[1]

    if pd.isnull(age):
        if pclass == 1:
            return round(df[df['Pclass']==1]['Age'].mean())
        elif pclass == 2:
```

```

    return round(df[df['Pclass']==2]['Age'].mean())
elif pclass == 3:
    return round(df[df['Pclass']==3]['Age'].mean())
else:
    return age
df['Age']=df[['Age', 'Pclass']].apply(fill_in_na_values, axis=1)

```

```

<ipython-input-27-d9409739061d>:2: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a futu
age = cols[0]
<ipython-input-27-d9409739061d>:3: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a futu
pclass = cols[1]

```

```
df.isnull().sum()
```

```

0
PassengerId  0
Pclass       0
Name         5
Sex          0
Age          0
SibSp        2
Parch        2
Ticket No.   2
Fare         2
Embarked     6
Survival     4

```

```
df.dropna(inplace=True)
```

```
df.isnull().sum()
```

```

0
PassengerId  0
Pclass       0
Name         0
Sex          0
Age          0
SibSp        0
Parch        0
Ticket No.   0
Fare         0
Embarked     0
Survival     0

```

```
df.head()
```

```

PassengerId  Pclass      Name  Sex  Age  SibSp  Parch  Ticket No.  Fare  Embarked  Survival
0           1      3.0  Braund, Mr. Owen Harris  male  22.0    1.0    0.0    A/5 21171  7.2500      S      0.0
2           3      3.0    Heikkinen, Miss. Laina  female  26.0    0.0    0.0  STON/O2.  7.9250      S      1.0
3           4      1.0  Futrelle, Mrs. Jacques Heath (Lily May Peel)  female  35.0    1.0    0.0    113803  53.1000      S      1.0

```

```
df.drop(['PassengerId', 'Name', 'Ticket No.'], axis=1, inplace=True)
```

```
df.head()
```

```
↗
```

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	Survival
0	3.0	male	22.0	1.0	0.0	7.2500	S	0.0
2	3.0	female	26.0	0.0	0.0	7.9250	S	1.0
3	1.0	female	35.0	1.0	0.0	53.1000	S	1.0
4	3.0	male	35.0	0.0	0.0	8.0500	S	0.0

```
df['Embarked'].unique()
```

```
↗ array(['S', 'C', 'Q'], dtype=object)
```

```
sex = pd.get_dummies(df['Sex'], drop_first=True).astype(int)
embarked = pd.get_dummies(df['Embarked'], drop_first=True).astype(int)
```

```
print(sex)
print(embarked)
```

```
↗
```

	male
0	1
2	0
3	0
4	1
6	1
..	...
885	1
886	0
887	0
888	1
889	1

[878 rows x 1 columns]

	Q	S
0	0	1
2	0	1
3	0	1
4	0	1
6	0	1
..	..	..
885	0	1
886	0	1
887	0	1
888	0	0
889	1	0

[878 rows x 2 columns]

Double-click (or enter) to edit

```
df = pd.concat([df, sex, embarked], axis=1)
```

Double-click (or enter) to edit

```
df.head()
```

```
↗
```

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	Survival	male	Q	S	male	Q	S
0	3.0	male	22.0	1.0	0.0	7.2500	S	0.0	True	False	True	1	0	1
2	3.0	female	26.0	0.0	0.0	7.9250	S	1.0	False	False	True	0	0	1
3	1.0	female	35.0	1.0	0.0	53.1000	S	1.0	False	False	True	0	0	1
4	3.0	male	35.0	0.0	0.0	8.0500	S	0.0	True	False	True	1	0	1

```
# Drop the first occurrences of 'male', 'Q', 'S' but keep the last three
df = df.drop(df.columns[list(df.columns).index('male'):list(df.columns).index('male')+3], axis=1)
```

```
df.head()
```

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	Survival
0	3.0	male	22.0	1.0	0.0	7.2500	S	0.0
2	3.0	female	26.0	0.0	0.0	7.9250	S	1.0
3	1.0	female	35.0	1.0	0.0	53.1000	S	1.0
4	3.0	male	35.0	0.0	0.0	8.0500	S	0.0

```
sex = pd.get_dummies(df['Sex'], drop_first=True).astype(int)
embarked = pd.get_dummies(df['Embarked'], drop_first=True).astype(int)
```

```
df = pd.concat([df, sex, embarked], axis=1)
```

```
df.head()
```

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	Survival	male	Q	S
0	3.0	male	22.0	1.0	0.0	7.2500	S	0.0	1	0	1
2	3.0	female	26.0	0.0	0.0	7.9250	S	1.0	0	0	1
3	1.0	female	35.0	1.0	0.0	53.1000	S	1.0	0	0	1
4	3.0	male	35.0	0.0	0.0	8.0500	S	0.0	1	0	1

```
df.drop(['Sex', 'Embarked'], axis=1, inplace=True)
```

```
df.head()
```

	Pclass	Age	SibSp	Parch	Fare	Survival	male	Q	S
0	3.0	22.0	1.0	0.0	7.2500	0.0	1	0	1
2	3.0	26.0	0.0	0.0	7.9250	1.0	0	0	1
3	1.0	35.0	1.0	0.0	53.1000	1.0	0	0	1
4	3.0	35.0	0.0	0.0	8.0500	0.0	1	0	1

```
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
```

```
x = df.drop('Survival', axis=1)
y = df['Survival']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1)
```

```
scaler = MinMaxScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
```

```
# Define X (features) and y (target)
X = df.drop('Survival', axis=1)
y = df['Survival']
```

```
# Split the last 140 rows as test set, and the rest as training set
X_train, X_test = X.iloc[:-140], X.iloc[-140:] # Last 140 rows for testing
y_train, y_test = y.iloc[:-140], y.iloc[-140:]
```

```
# Normalize features using MinMaxScaler
scaler = MinMaxScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test) # Use same scaler for test data
```

```
from sklearn.svm import SVC
svm= SVC()
```

```
print(classification_report(y_test, grid_predictions))
print(confusion_matrix(y_test, grid_predictions))
```



```

↗
precision    recall  f1-score   support

0.0         0.81    0.97    0.88        90
1.0         0.91    0.58    0.71        50

accuracy
macro avg    0.86    0.77    0.79    140
weighted avg 0.84    0.83    0.82    140

[[87  3]
 [21 29]]

```

Start coding or [generate](#) with AI.

```

from sklearn.linear_model import LogisticRegression
lm = LogisticRegression()
lm.fit(X_train, y_train)
lm_prediction = lm.predict(X_test)

```

```

print(classification_report(y_test, lm_prediction))
print(confusion_matrix(y_test, lm_prediction))

```

```

↗
precision    recall  f1-score   support

0.0         0.83    0.89    0.86        90
1.0         0.77    0.68    0.72        50

accuracy
macro avg    0.80    0.78    0.79    140
weighted avg 0.81    0.81    0.81    140

[[80 10]
 [16 34]]

```

```

from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier()
knn.fit(X_train, y_train)
knn_prediction = knn.predict(X_test)

```

```

print(classification_report(y_test, knn_prediction))
print(confusion_matrix(y_test, knn_prediction))

```

```

↗
precision    recall  f1-score   support

0.0         0.83    0.88    0.85        90
1.0         0.76    0.68    0.72        50

accuracy
macro avg    0.79    0.78    0.78    140
weighted avg 0.80    0.81    0.80    140

[[79 11]
 [16 34]]

```

```
import numpy as np
```

```

from sklearn.neighbors import KNeighborsClassifier
import numpy as np

```

```
error_list = []
```

```

# Loop over different values of k (from 1 to 39)
for i in range(1, 40):
    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(X_train, y_train)
    prediction_i = knn.predict(X_test)

```

```

# Append the error rate for the current k (n_neighbors=i)
error_list.append(np.mean(prediction_i != y_test))

```

```

# Optionally, you could plot or analyze the error_list to see the optimal k
print(error_list)
import matplotlib.pyplot as plt

```


```

# Plotting the error rate vs. number of neighbors
plt.plot(range(1, 40), error_list, marker='o')
plt.title('Error rate vs. Number of Neighbors')
plt.xlabel('Number of Neighbors (k)')
plt.ylabel('Error Rate')
plt.show()

```



Error rate vs. Number of Neighbors



```
np.argmin(error_list)
```



29

```
error_list[6]
```



0.18571428571428572

```
knn = KNeighborsClassifier()
knn.fit(X_train, y_train)
knn_prediction = knn.predict(X_test)
print(classification_report(y_test, knn_prediction))
print(confusion_matrix(y_test, knn_prediction))
```



	precision	recall	f1-score	support
0.0	0.83	0.88	0.85	90
1.0	0.76	0.68	0.72	50
accuracy			0.81	140
macro avg	0.79	0.78	0.78	140
weighted avg	0.80	0.81	0.80	140

```
[[79 11]
 [16 34]]
```

```
from sklearn.tree import DecisionTreeClassifier
dtree = DecisionTreeClassifier()
dtree.fit(X_train, y_train)
dtree_prediction = dtree.predict(X_test)
print(classification_report(y_test, dtree_prediction))
print(confusion_matrix(y_test, dtree_prediction))
```



	precision	recall	f1-score	support
0.0	0.90	0.87	0.88	90
1.0	0.77	0.82	0.80	50
accuracy			0.85	140
macro avg	0.84	0.84	0.84	140
weighted avg	0.85	0.85	0.85	140

```
[[78 12]
 [ 9 41]]
```

```
from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(X_train, y_train)
rfc_prediction = rfc.predict(X_test)
print(classification_report(y_test, rfc_prediction))
```



	precision	recall	f1-score	support
0.0	0.90	0.90	0.90	90